

# CS 744 : DECS Project : Phase 1

## Key-Value Server With Database and LRU Cache Integration

[Github Repository](#)

## System Overview

This project implements a **client–server key–value storage system** that provides both **fast in-memory access** and **persistent database storage**. The system is designed to efficiently handle frequent read and write operations using an **LRU (Least Recently Used) cache** for hot data and **PostgreSQL** for long-term persistence.

---

## System Architecture

The architecture follows a modular design consisting of three core components:


### 1. Client

The client module provides the interface for users or external applications to send key–value operations (e.g., **GET**, **PUT**, **DELETE**) to the server. It communicates over TCP sockets using a simple request–response protocol.

### 2. Server Core

The server listens for client connections, parses incoming commands, and coordinates data access between the cache and the database.

- On a **read request**, it first queries the **LRU cache**. If the key exists in the cache, the value is returned immediately.
- On a **cache miss**, it fetches the value from **PostgreSQL**, updates the cache, and responds to the client.
- On a **write request**, it updates both the cache and the database to maintain consistency.



### 3. LRU Cache Module

This in-memory cache is implemented using a combination of a **hash map** and a **doubly-linked list**, enabling O(1) access, insertion, and eviction operations. When the cache reaches capacity, the least recently used item is evicted automatically.

### 4. Database Layer (PostgreSQL)

Acts as the persistent storage backend. All key–value pairs are stored in a `kv_table` with schema `(key TEXT PRIMARY KEY, value TEXT)`. This ensures data durability across restarts and enables recovery from cache losses.

---

## Data Flow Summary

1. **Client → Server:** Sends request (e.g., `PUT key value` or `GET key`).
  2. **Server → Cache:** Checks or updates the LRU cache.
  3. **Server → Database:** Syncs persistent data if necessary.
  4. **Server → Client:** Sends back the response.
- 

## Key Features

- **Hybrid storage:** Combines in-memory speed with persistent reliability.
- **Scalable and modular:** Components can be extended or replaced independently.
- **Efficient caching:** LRU policy ensures frequently accessed data remains hot.
- **Crash recovery:** Persistent PostgreSQL storage guarantees data integrity.